

# **Lending Club Loan Default Detection**

# CONTENT

## **01 Project Overview**

- ❖ Project Introduction
- ❖ Data Overview

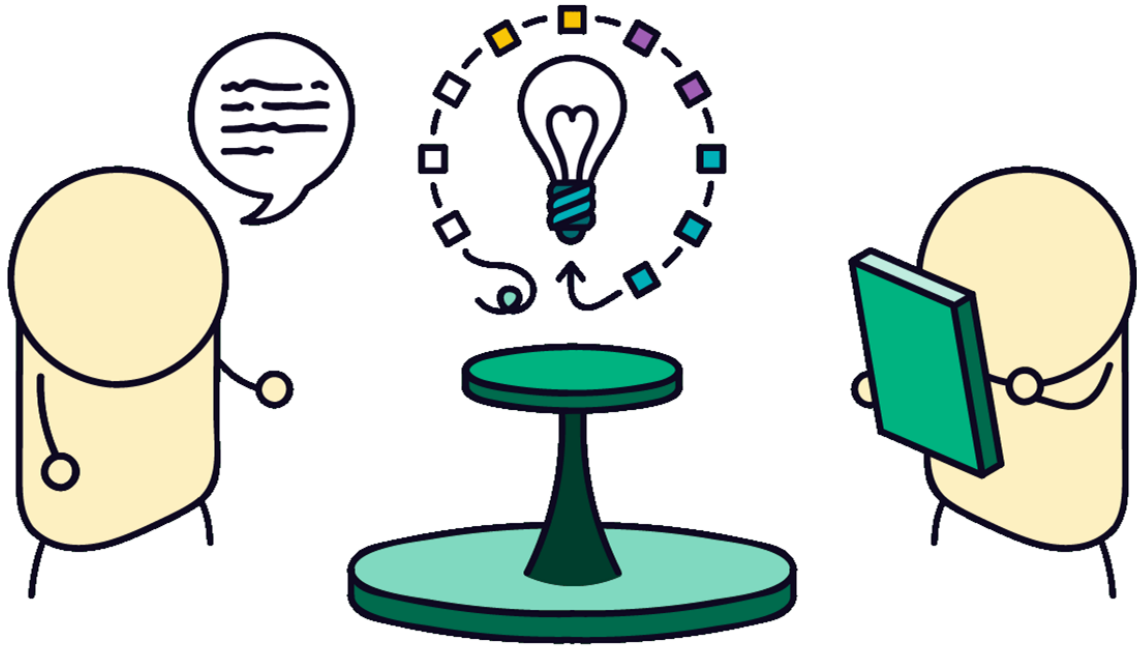
## **02 Data Preprocessing**

## **03 Loan Default Prediction (modeling)**

- ❖ Supervised Learning: Probability of Default Estimation
- ❖ Self-supervised Learning: Anomaly Detection

## **04 Model Performance Comparison**

## **05 Conclusions**



# 1 Project Overview

Project Introduction

Problem Description

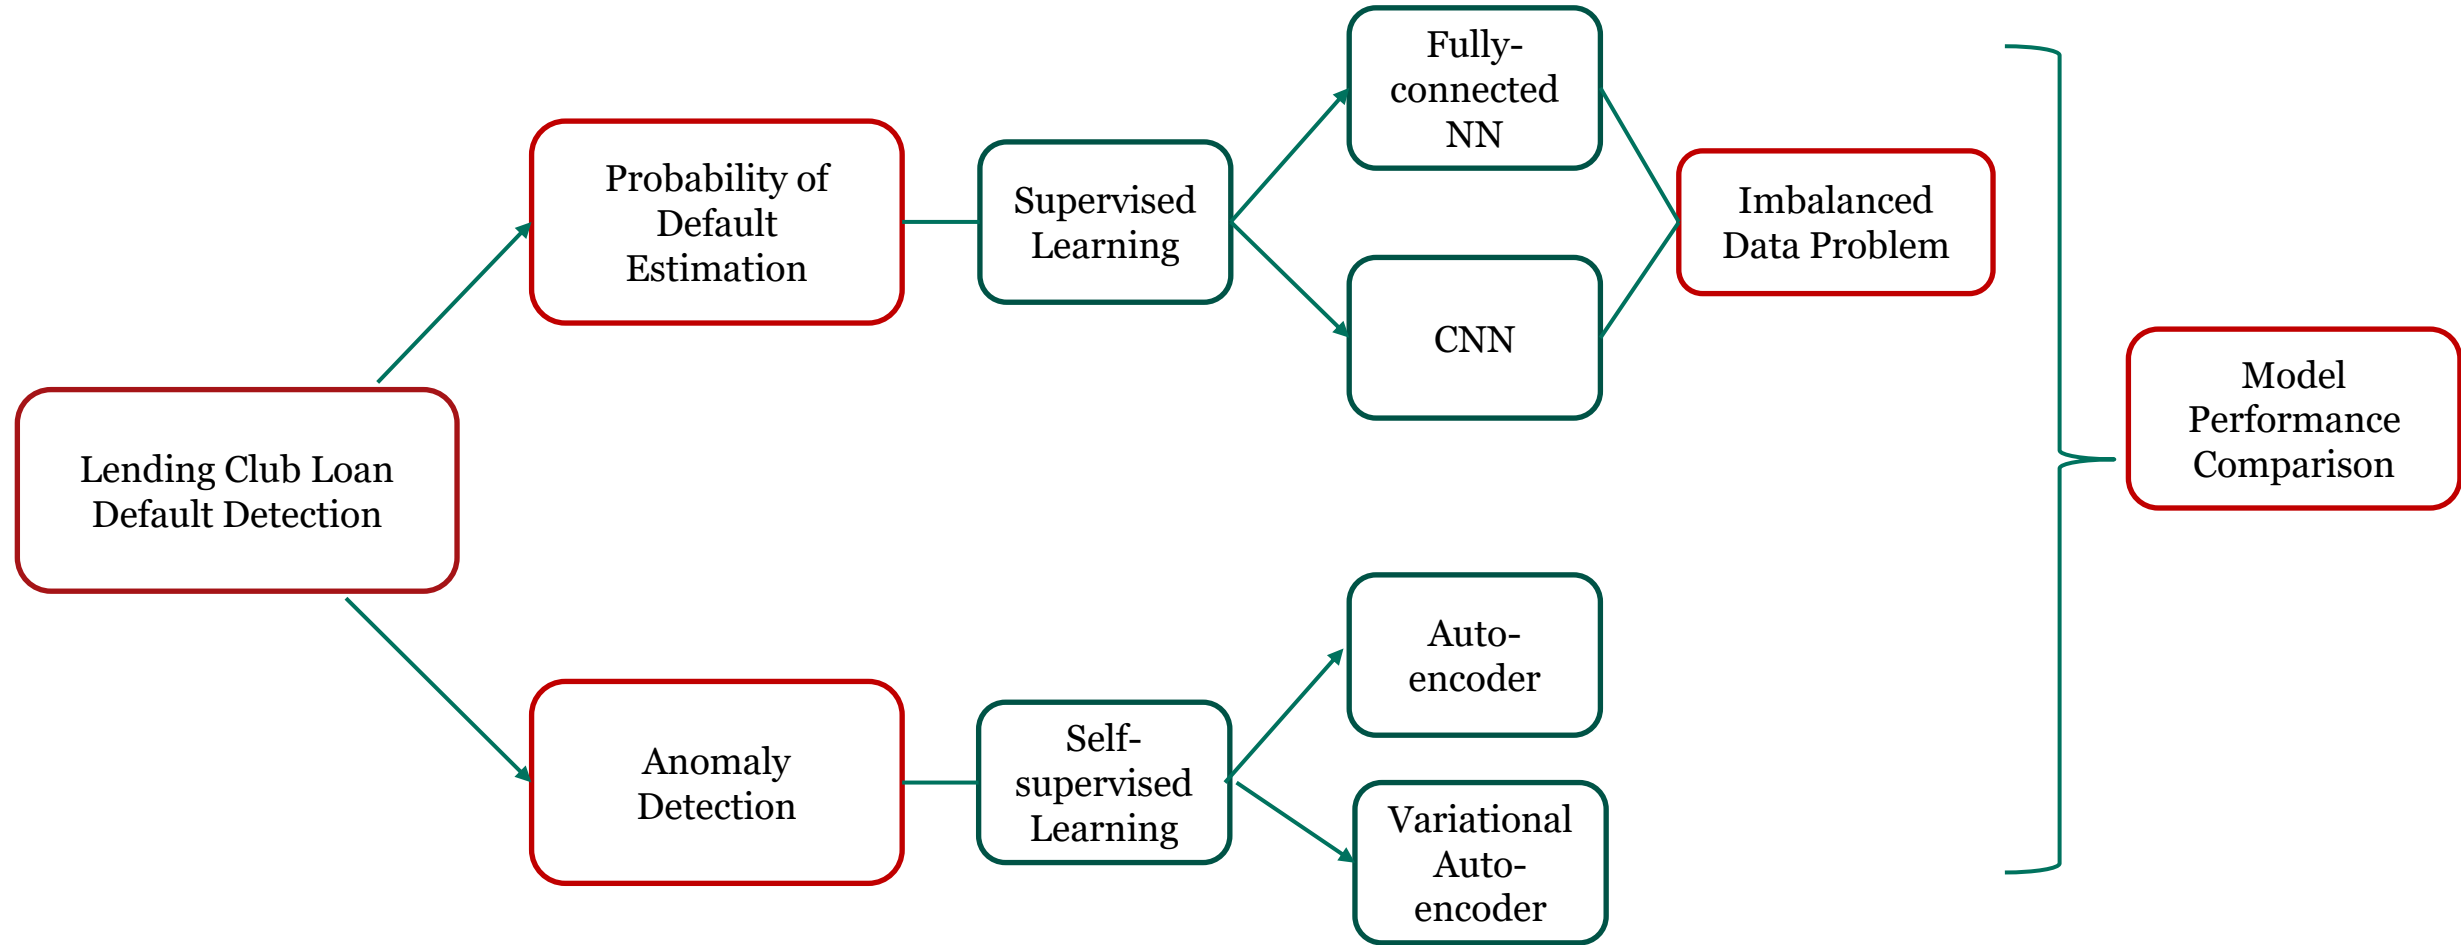
Data Overview

# 1.1 Project Introduction

- Peer to peer lending & financial inclusion
- Objectives: Loan repayment defaulter detection



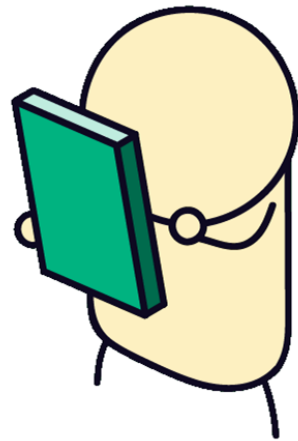
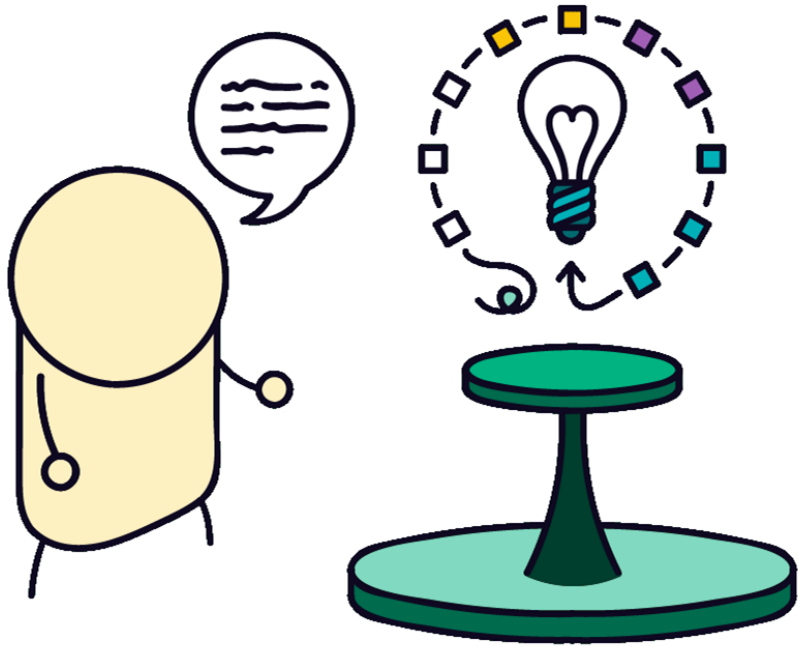
## 1.2 Problems Description



# 1.3 Database Introduction

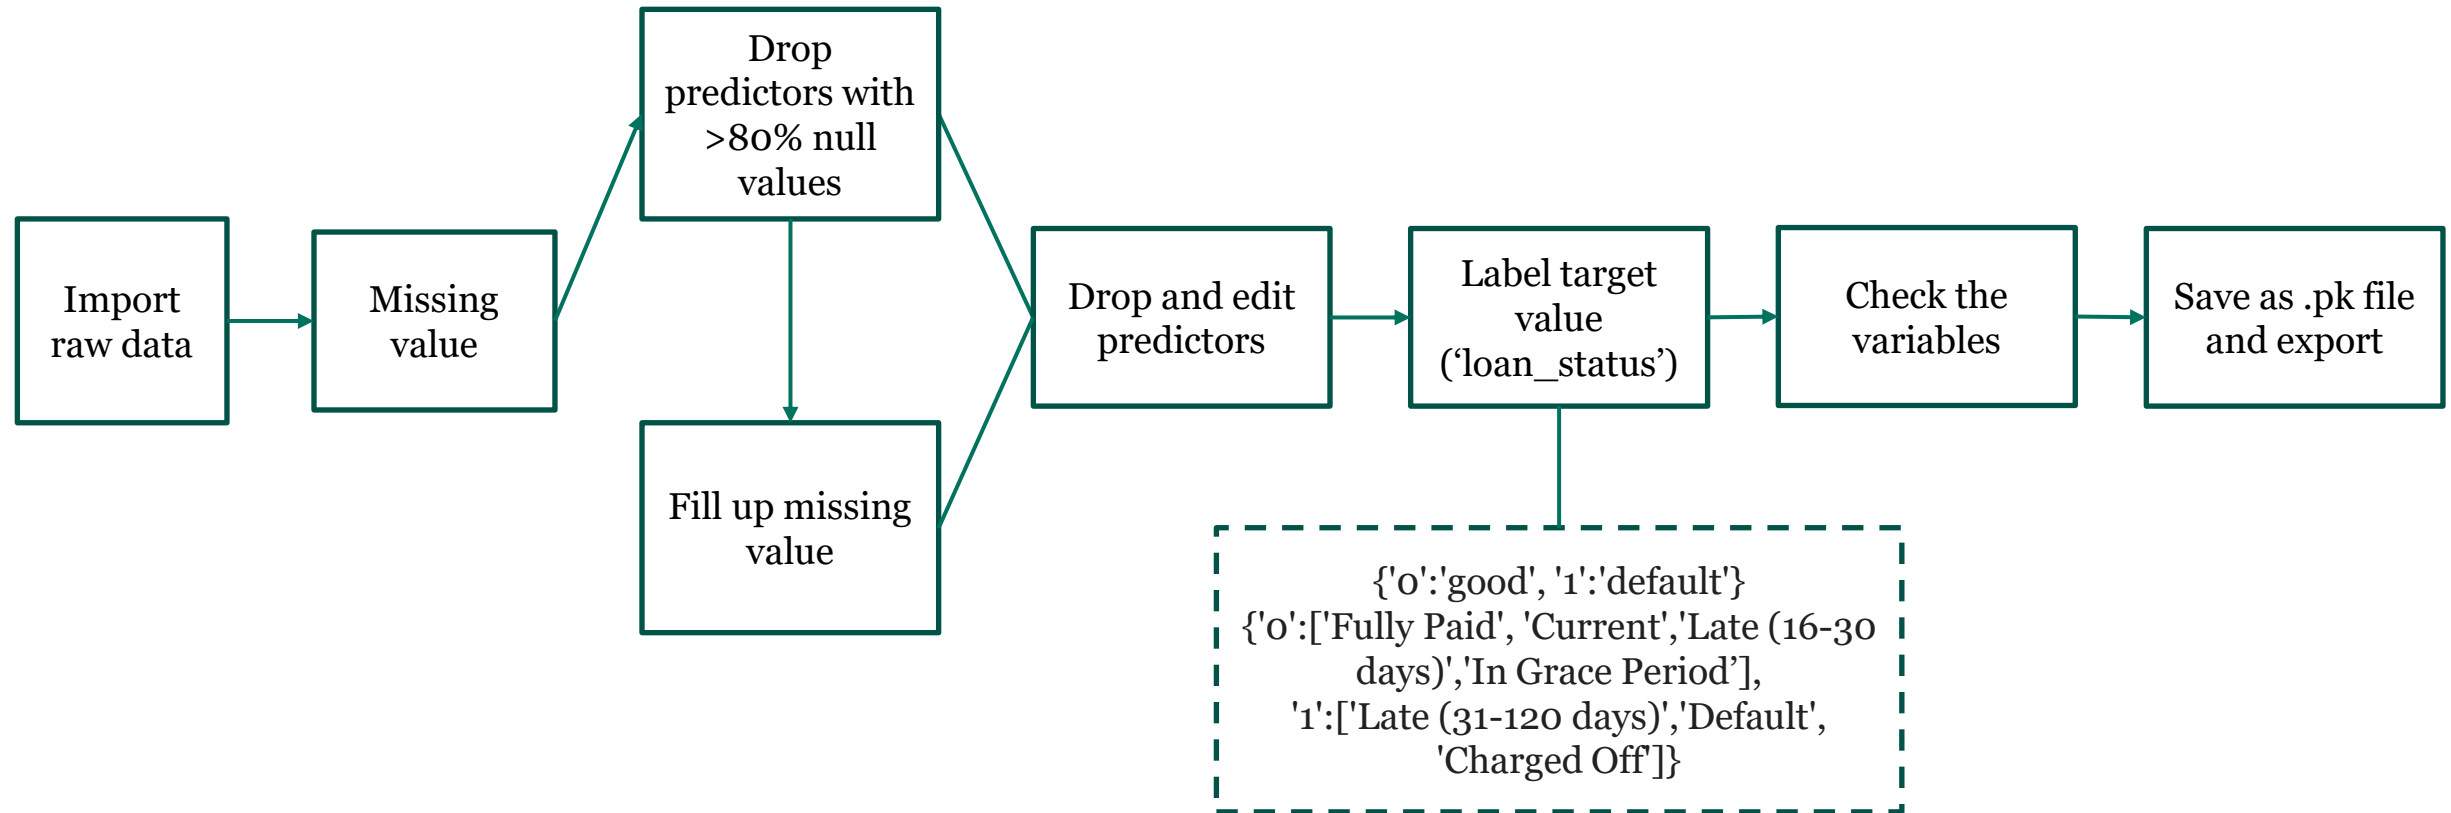
LendingClub Issued Loans Data: <https://www.kaggle.com/husainsb/lendingclub-issued-loans>

- Training Data set: lc\_loan.csv  
Contains loans issued from 2007-2015  
74 columns
- Test Data set: lc\_2016\_2017.csv  
Contains loans issued from 2016-2017  
72 columns (missing the columns of 'open\_il\_6m', 'url')



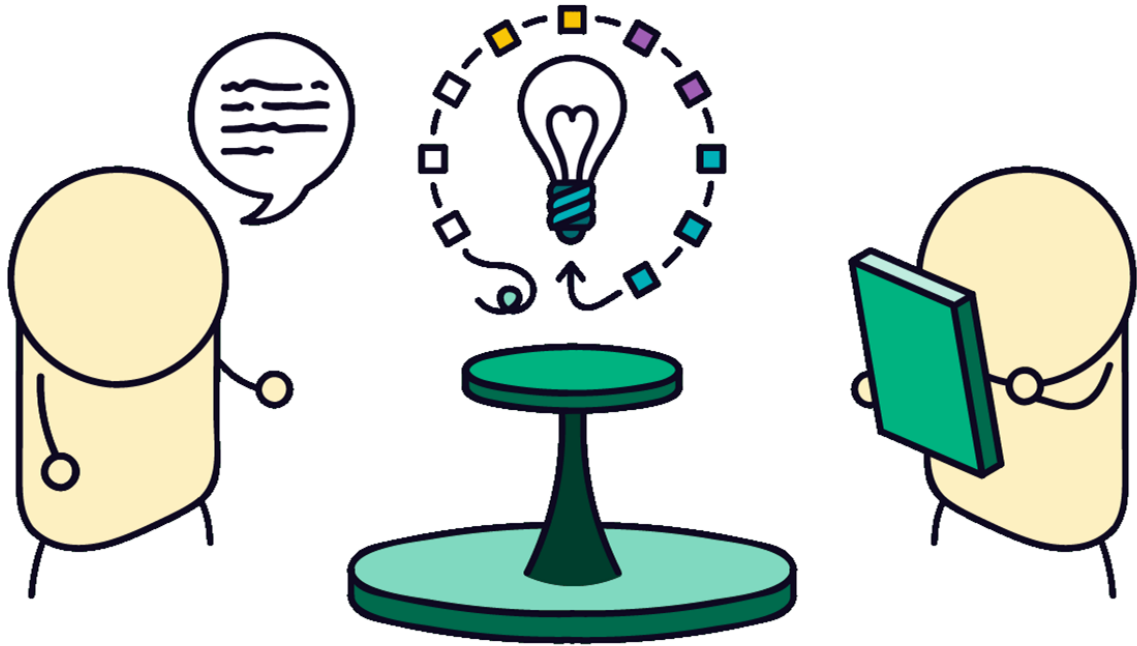
## 2 Data preprocessing

# Data Preprocessing



- Predictors (X): from 73 to 41
- Target Variables (Y): {'0': 'good', '1': 'default'}





## 3 Modeling -- Loan Default Prediction

- 1) Supervised Learning
- 2) Unsupervised Learning

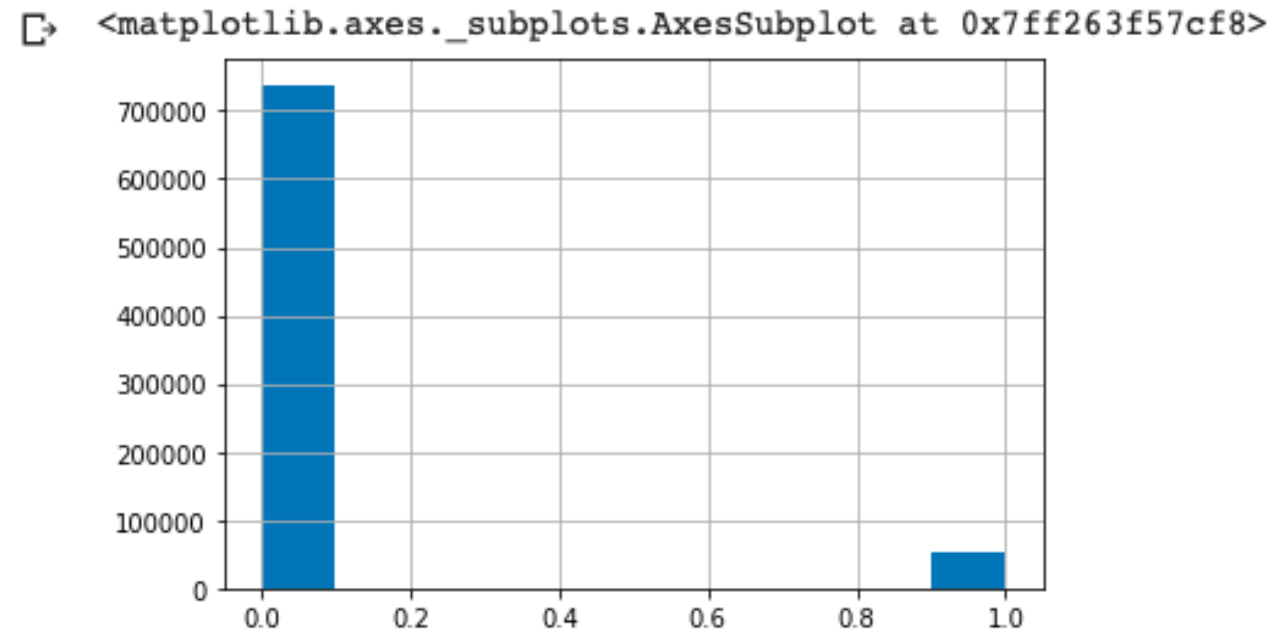
## 3.1 Probability of Default Estimation

-- Fully-connected Neural Network

- **Supervised Learning**
- **Fully-connected Neural Network**
- **Undersampling**
- **Oversampling**

# Problem of Imbalanced Classification

- Target column -- “**Loan\_Status**” 0 is good 1 is default
- The number of default transactions in training data is **52937** and number of transactions which do not default is **738087**;
- The number of default transactions in testing data is **52564** and number of transactions which do not default is **706270**;



# 1 Basic Model -- Fully-connected Neural Network

## 1.1 Modeling

- Standardization
- Dropout
- Regularization
- Change architecture
- Early Stopping

Model: "sequential\_15"

Layer (type)	Output Shape	Param #
dense_58 (Dense)	(None, 128)	5376
dense_59 (Dense)	(None, 64)	8256
dropout_38 (Dropout)	(None, 64)	0
dense_60 (Dense)	(None, 32)	2080
dropout_39 (Dropout)	(None, 32)	0
dense_61 (Dense)	(None, 16)	528
dense_62 (Dense)	(None, 1)	17

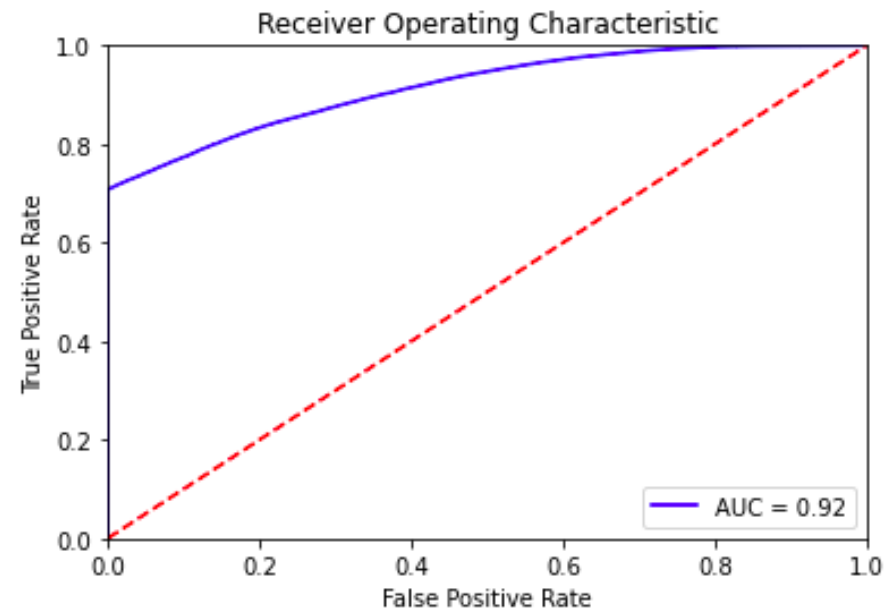
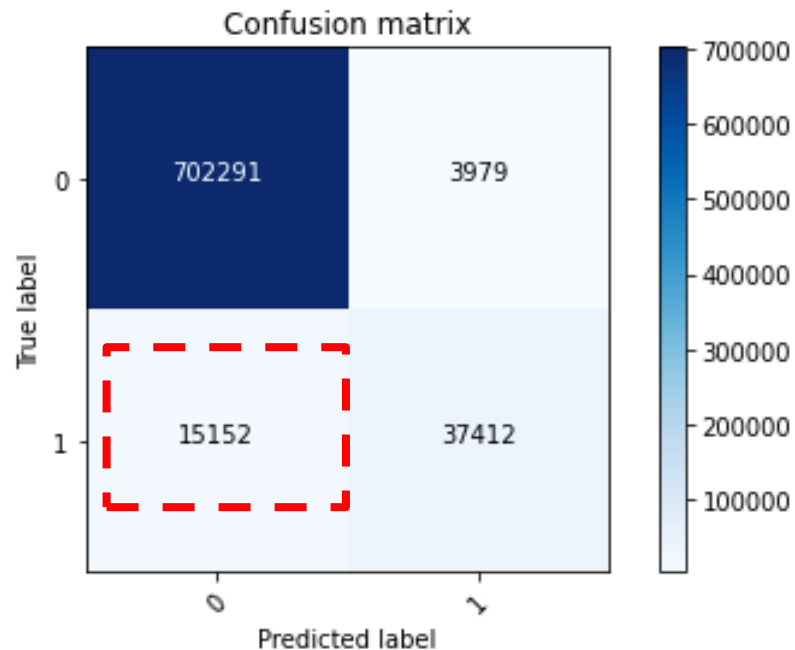
Total params: 16,257  
Trainable params: 16,257  
Non-trainable params: 0

# 1 Basic Model -- Fully-connected Neural Network

## 1.2 Evaluation

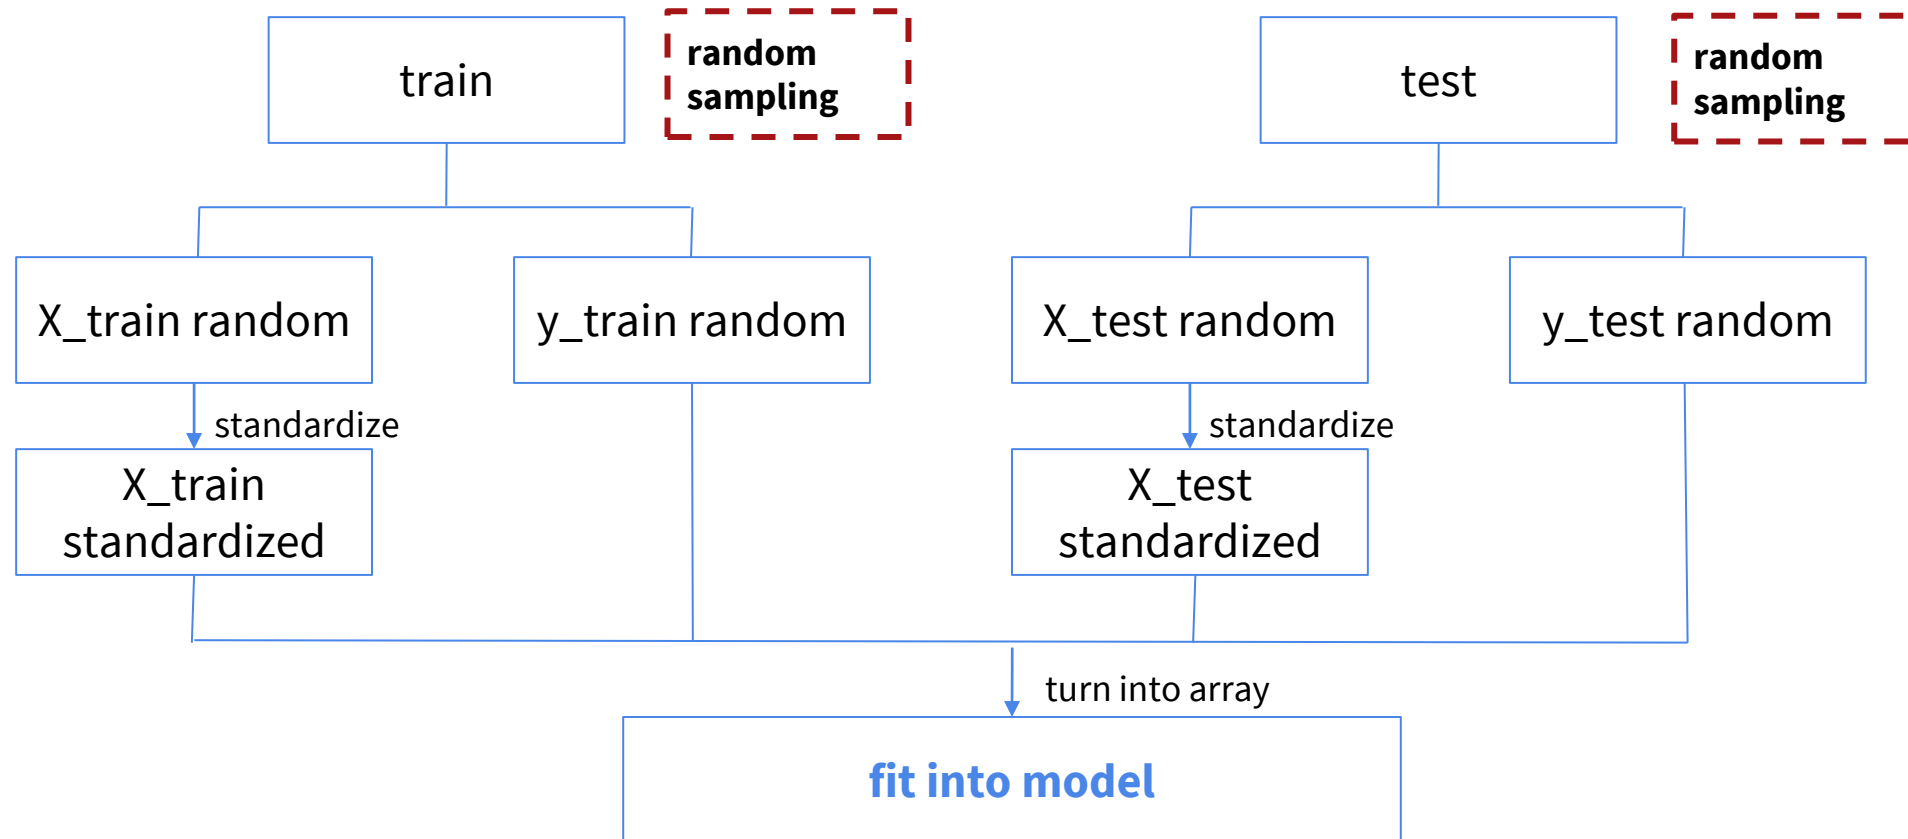
```
print("The testing accuracy is:", model.evaluate(X_test, y_test))
```

23714/23714 [=====] - 49s 2ms/step - loss: 0.1211 - accuracy: 0.9748  
The testing accuracy is: [0.12106689065694809, 0.9747889637947083]



## 2 Undersampling

- **reducing** the data by eliminating examples belonging to the **majority class**



# 2 Undersampling

## 2.1 Modeling

- Standardization
- Regularization
- Change architecture
- Early Stopping

Model: "sequential\_16"

Layer (type)	Output Shape	Param #
dense_63 (Dense)	(None, 128)	5376
dense_64 (Dense)	(None, 64)	8256
dense_65 (Dense)	(None, 32)	2080
dense_66 (Dense)	(None, 16)	528
dense_67 (Dense)	(None, 1)	17

Total params: 16,257

Trainable params: 16,257

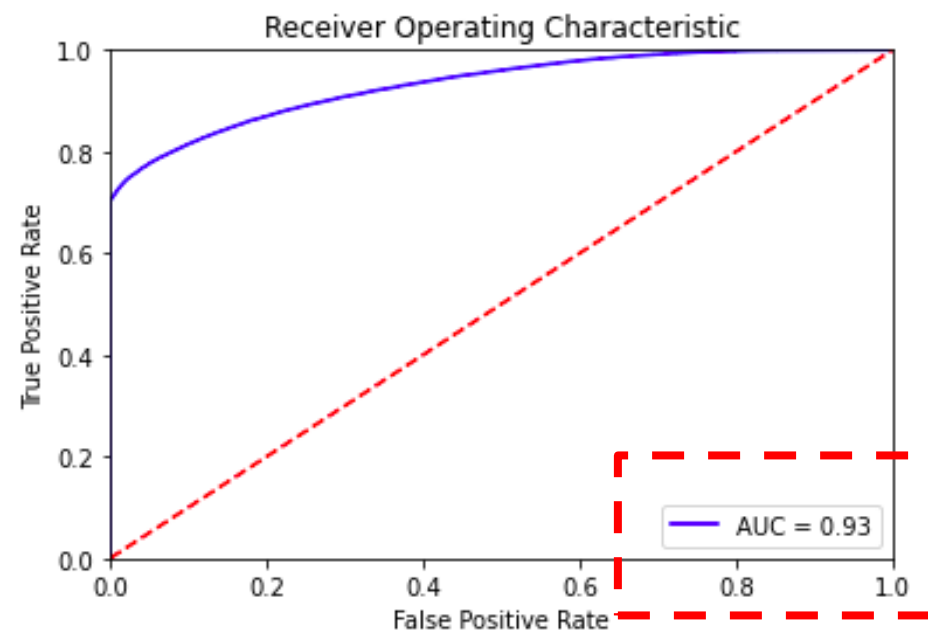
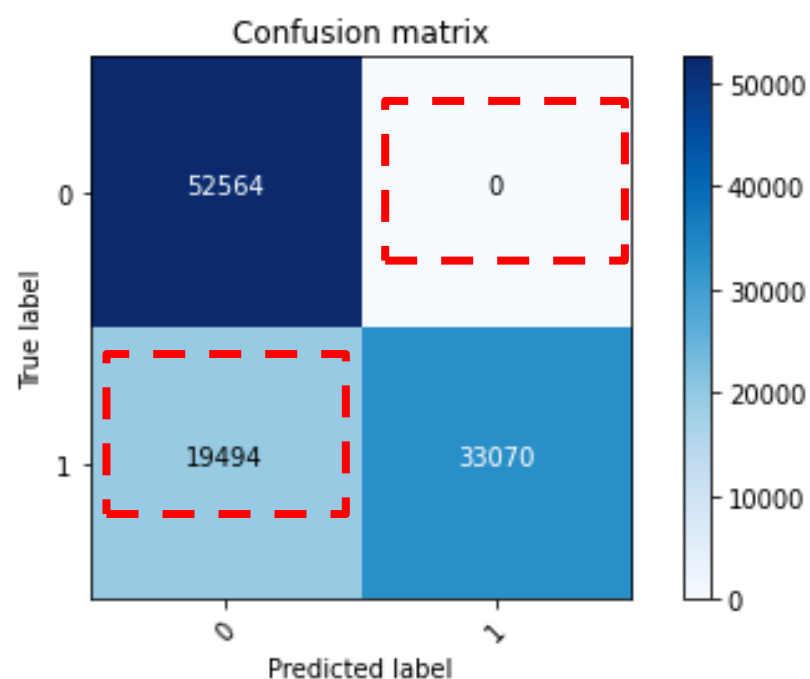
Non-trainable params: 0

# 2 Undersampling

## 2.2 Evaluation

```
print("The testing accuracy is:", under_model.evaluate(under_X_test, under_y_test))
```

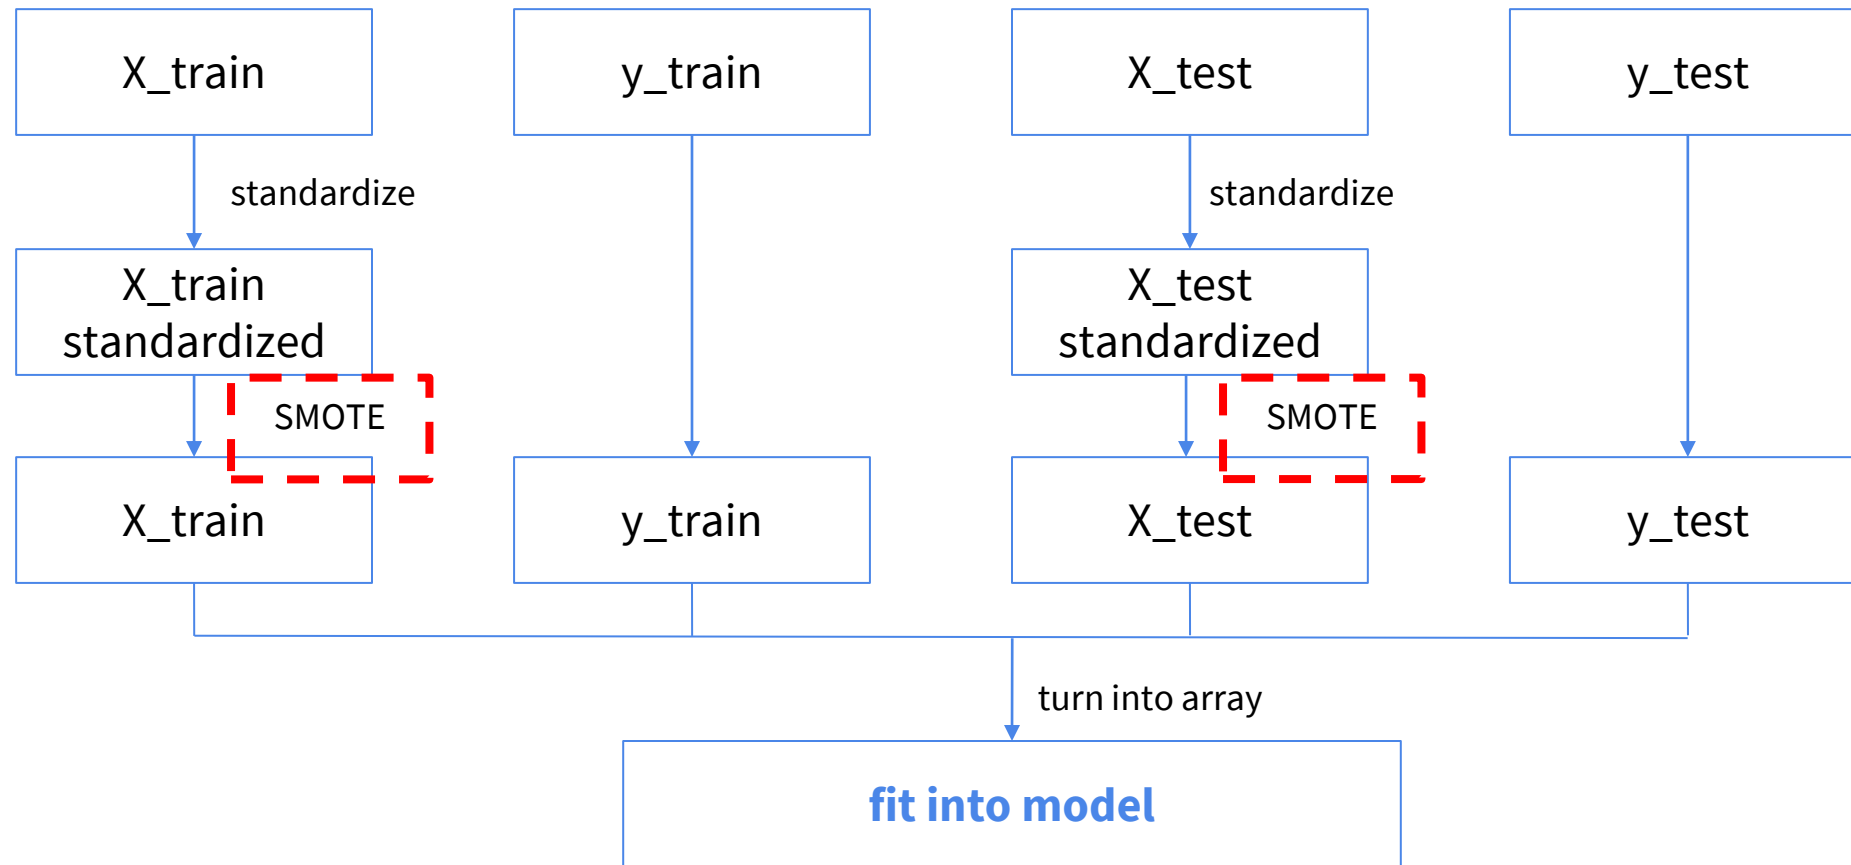
3286/3286 [=====] - 5s 2ms/step - loss: 0.2901 - accuracy: 0.8642  
The testing accuracy is: [0.29013320803642273, 0.8641656041145325]





# 3 Oversampling

- selecting examples that are close in the feature space, drawing a line between the examples in the feature space and drawing a new sample at a point along that line -- increase the percentage of minority class



# 3 Oversampling

## 3.1 Modeling

- Standardization
- Dropout
- Regularization
- Change architecture
- Early Stopping

Model: "sequential\_18"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_73 (Dense)	(None, 64)	2688
dropout_44 (Dropout)	(None, 64)	0
dense_74 (Dense)	(None, 32)	2080
dropout_45 (Dropout)	(None, 32)	0
dense_75 (Dense)	(None, 32)	1056
dropout_46 (Dropout)	(None, 32)	0
dense_76 (Dense)	(None, 16)	528
dropout_47 (Dropout)	(None, 16)	0
dense_77 (Dense)	(None, 1)	17
=====	=====	=====

Total params: 6,369

Trainable params: 6,369

Non-trainable params: 0

# 3 Oversampling

## 3.2 Evaluation

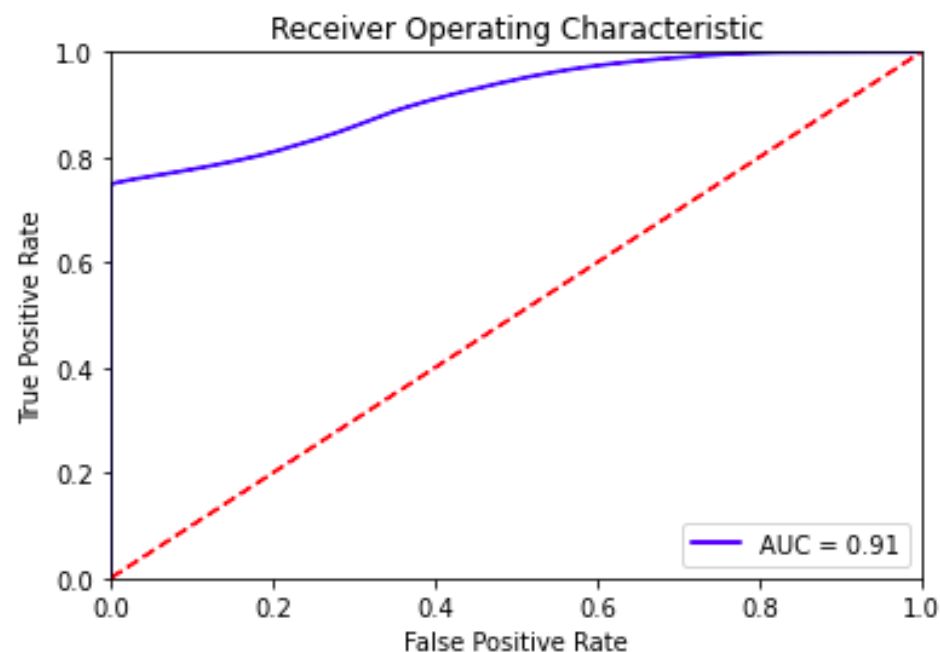
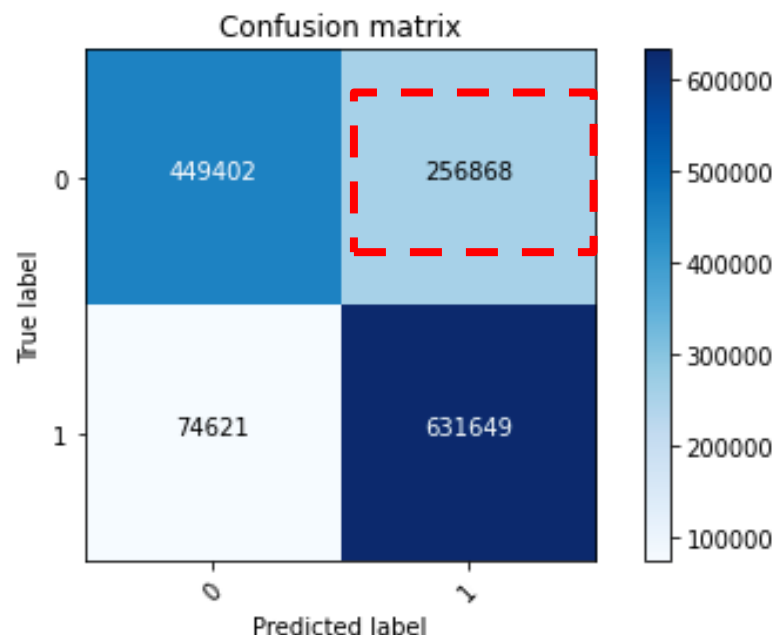
```
print("The testing accuracy is:", over_model.evaluate(X_test_resample, y_test_resample))
```

```
44142/44142 [=====] - 75s 2ms/step - loss: 0.6918 - accuracy: 0.7653  
The testing accuracy is: [0.6918033957481384, 0.7653241753578186]
```

Confusion matrix, without normalization

```
[[449402 256868]
```

```
 [ 74621 631649]]
```



# 4 Conclusion

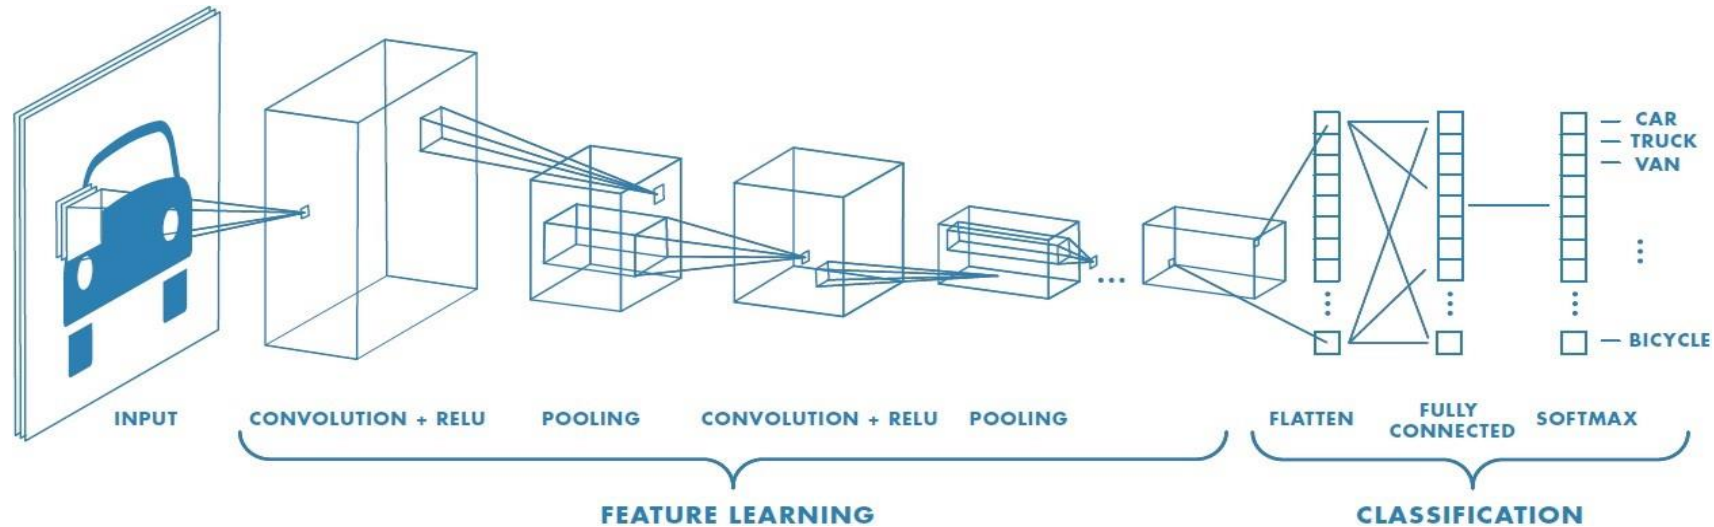
- Utilizing dropout, regularization, change architecture, and early stopping, overfitting problem can be solved well
- Fully connected neural model has a good accuracy
- Undersampling would ease imbalance problem
- Oversampling might not be appropriate to solve the problem in this dataset

## 3.2 Probability of Default Estimation

### -- Convolutional Neural Network

- **Convolutional Neural Network**
- **Undersampling**
- **Oversampling**

# Convolutional Neural Network Overview



- One of the main categories to do images recognition, images classifications
- Also can be used for non-image data set by reshaping the data with the input dimension = 1

# 1 Modeling

- Convolutional layer
- Dropout
- MaxPooling
- Flatten
- Early Stopping

Model: "model\_conv1D"

Layer (type)	Output Shape	Param #
Conv1D_1 (Conv1D)	(None, 35, 64)	512
dropout (Dropout)	(None, 35, 64)	0
Conv1D_2 (Conv1D)	(None, 33, 32)	6176
Conv1D_3 (Conv1D)	(None, 32, 16)	1040
MaxPooling1D (MaxPooling1D)	(None, 16, 16)	0
flatten (Flatten)	(None, 256)	0
Dense_1 (Dense)	(None, 32)	8224
dense (Dense)	(None, 1)	33
Total params: 15,985		
Trainable params: 15,985		
Non-trainable params: 0		

## 2 Evaluate

```
history = model_conv1D.fit(x_train_resaped, y_train.values, epochs=10,  
                           validation_data = (x_validation_resaped, y_validation.values), callbacks = [Es])
```

```
Epoch 1/10  
24720/24720 [=====] - 196s 8ms/step - loss: 7.3251 - acc: 0.9285 - val_loss: 0.1297 - val_acc: 0.9742  
Epoch 2/10  
24720/24720 [=====] - 182s 7ms/step - loss: 0.1418 - acc: 0.9724 - val_loss: 0.1073 - val_acc: 0.9782  
Epoch 3/10  
24720/24720 [=====] - 178s 7ms/step - loss: 0.1270 - acc: 0.9754 - val_loss: 0.1101 - val_acc: 0.9784  
Epoch 4/10  
24720/24720 [=====] - 175s 7ms/step - loss: 0.1244 - acc: 0.9759 - val_loss: 0.0924 - val_acc: 0.9833  
Epoch 5/10  
24720/24720 [=====] - 174s 7ms/step - loss: 0.0994 - acc: 0.9790 - val_loss: 0.0859 - val_acc: 0.9821  
Epoch 6/10  
24720/24720 [=====] - 180s 7ms/step - loss: 0.1055 - acc: 0.9788 - val_loss: 0.1079 - val_acc: 0.9772  
Epoch 7/10  
24720/24720 [=====] - 177s 7ms/step - loss: 0.1118 - acc: 0.9780 - val_loss: 0.0892 - val_acc: 0.9816
```

```
print("The testing accuracy is:", model_conv1D.evaluate(x_test_resaped, y_test.values))
```

```
23714/23714 [=====] - 48s 2ms/step - loss: 0.0999 - acc: 0.9791  
The testing accuracy is: [0.09991364181041718, 0.9791087508201599]
```

- High accuracy with no overfitting
- But, data imbalance might lead to the “perfect” outcome  
→ **Need oversample/undersample**



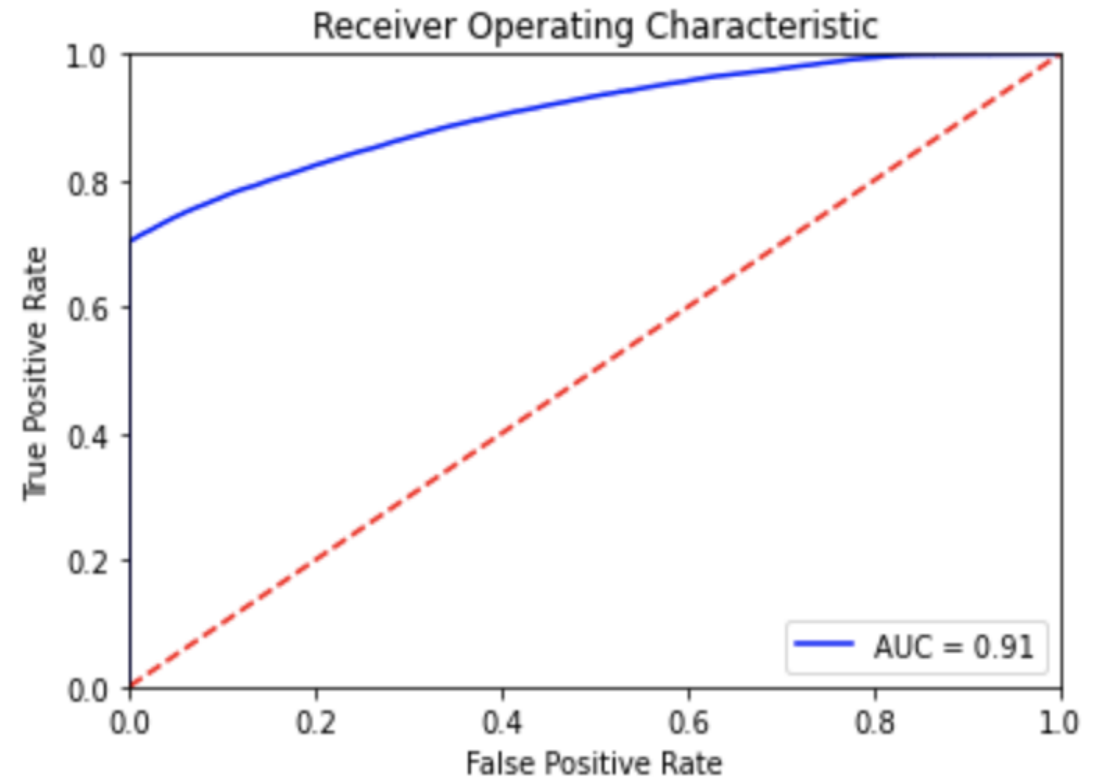
# 3 Undersampling

```
model_conv1D = build_conv1D_model()
history = model_conv1D.fit(x_train_bal_resaped, y_train_bal, epochs=10,
                           validation_data = (x_validation_bal_resaped, y_validation_bal),ca
```

```
Epoch 1/10
3309/3309 [=====] - 26s 8ms/step - loss: 114.8748 - acc: 0.64
Epoch 2/10
3309/3309 [=====] - 25s 8ms/step - loss: 0.9238 - acc: 0.7129
Epoch 3/10
3309/3309 [=====] - 25s 8ms/step - loss: 0.5286 - acc: 0.7759
Epoch 4/10
3309/3309 [=====] - 26s 8ms/step - loss: 0.4402 - acc: 0.8221
Epoch 5/10
3309/3309 [=====] - 26s 8ms/step - loss: 0.3337 - acc: 0.8616
Epoch 6/10
3309/3309 [=====] - 25s 8ms/step - loss: 0.3282 - acc: 0.8605
Epoch 7/10
3309/3309 [=====] - 25s 8ms/step - loss: 0.3198 - acc: 0.8632
Epoch 8/10
3309/3309 [=====] - 25s 8ms/step - loss: 0.3048 - acc: 0.8699
Epoch 9/10
3309/3309 [=====] - 26s 8ms/step - loss: 0.3101 - acc: 0.8675
```

```
print("The testing accuracy is:", model_conv1D.evaluate(x_test_bal_resaped, y_test_ba
```

```
3286/3286 [=====] - 8s 2ms/step - loss: 0.3205 - acc: 0.8516
The testing accuracy is: [0.3204999566078186, 0.8515619039535522]
```



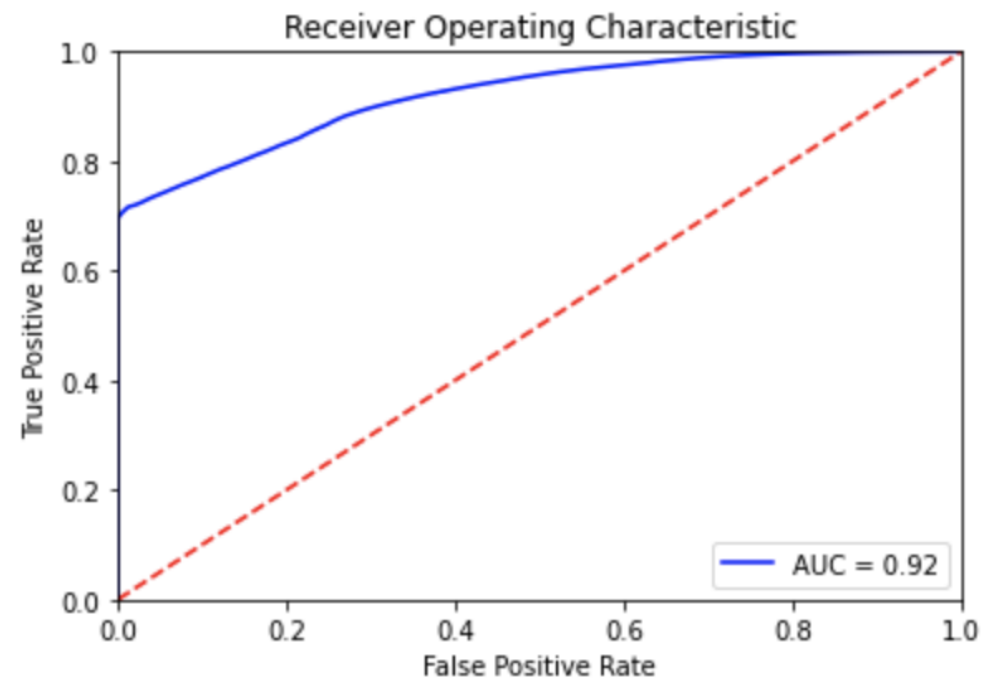
# 4 Oversampling

```
model_conv1D = build_conv1D_model()
history = model_conv1D.fit(x_train_bal_resaped, y_train_bal, epochs=10,
                           validation_data = (x_validation_bal_resaped, y_validation_bal), callbacks = [Es
```

```
Epoch 1/10
46131/46131 [=====] - 315s 7ms/step - loss: 14.0595 - acc: 0.7959 - val_loss: 0.2897
Epoch 2/10
46131/46131 [=====] - 337s 7ms/step - loss: 0.2897 - acc: 0.8715 - val_loss: 0.2828
Epoch 3/10
46131/46131 [=====] - 329s 7ms/step - loss: 0.2828 - acc: 0.8747 - val_loss: 0.2805
Epoch 4/10
46131/46131 [=====] - 329s 7ms/step - loss: 0.2805 - acc: 0.8733 - val_loss: 0.2730
Epoch 5/10
46131/46131 [=====] - 329s 7ms/step - loss: 0.2730 - acc: 0.8769 - val_loss: 0.2764
Epoch 6/10
46131/46131 [=====] - 330s 7ms/step - loss: 0.2764 - acc: 0.8745 - val_loss: 0.2752
Epoch 7/10
46131/46131 [=====] - 331s 7ms/step - loss: 0.2752 - acc: 0.8762 - val_loss: 0.2826
Epoch 8/10
46131/46131 [=====] - 329s 7ms/step - loss: 0.2826 - acc: 0.8763 - val_loss: 0.2826
```

```
print("The testing accuracy is:", model_conv1D.evaluate(x_test_bal_resaped, y_test_bal))
```

```
44142/44142 [=====] - 82s 2ms/step - loss: 0.3100 - acc: 0.8485
The testing accuracy is: [0.3100346624851227, 0.8484963178634644]
```



## 5 Conclusion

- Though CNN result for original data works pretty well, it cannot be taken into account due to imbalance data set.
- Oversampled data performs better with the CNN model in both accuracy and AUC/ROC curve.

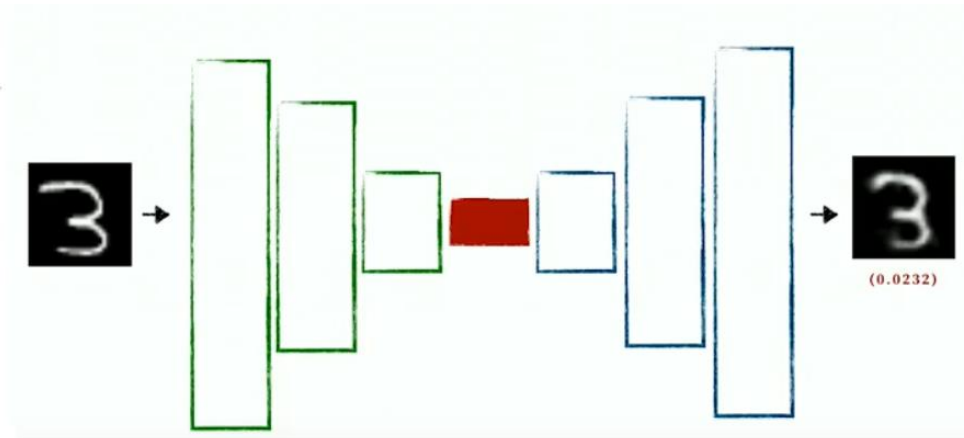
## 3.3 Anomaly Detection

-- self-supervised Learning

- **Anomaly Detection**
- **Autoencoder**
- **Variational autoencoder**

# 1 Anomaly Detection

## Self-supervised Learning -- Autoencoder/VAE

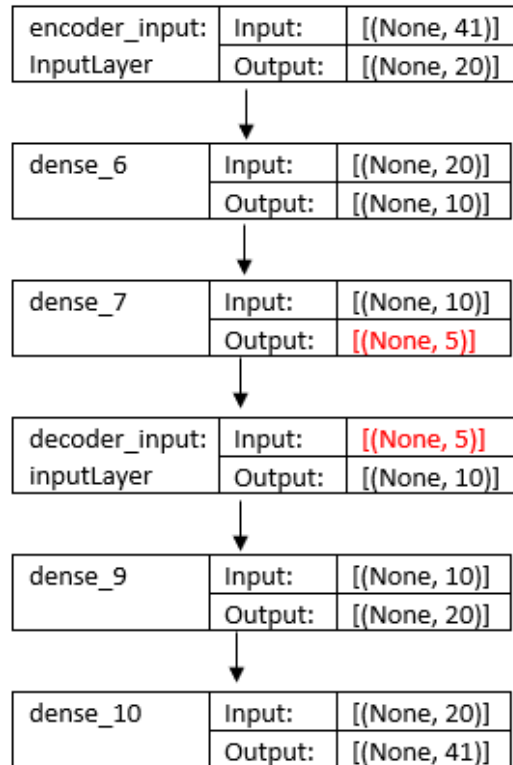


- Train model using **normal training data set**
- Find the threshold of the anomaly using reconstruction error (mean-squared-error)
- If reconstruction error  $>$  threshold, it's an anomaly

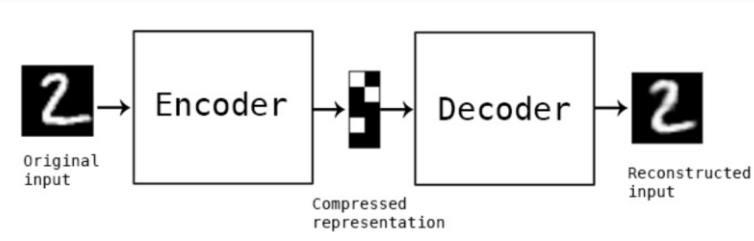
# 2 Autoencoder

## Modeling

Basic Model



val\_loss : 0.0032  
val\_accuracy: 0.5886

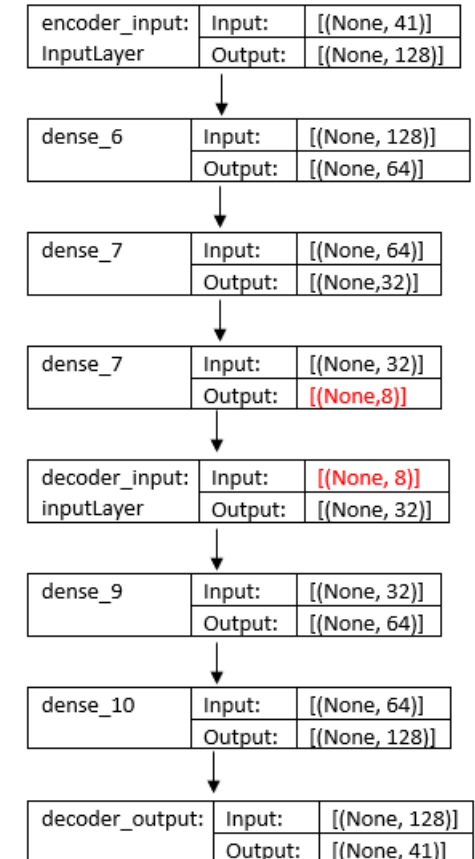


Parameter: loss='MSE', optimizer='adam'

- Change architectures
  - Bottleneck
  - Hidden layers
- Parameters
  - Loss function
- Early Stopping

Model Improvement

Improved Model

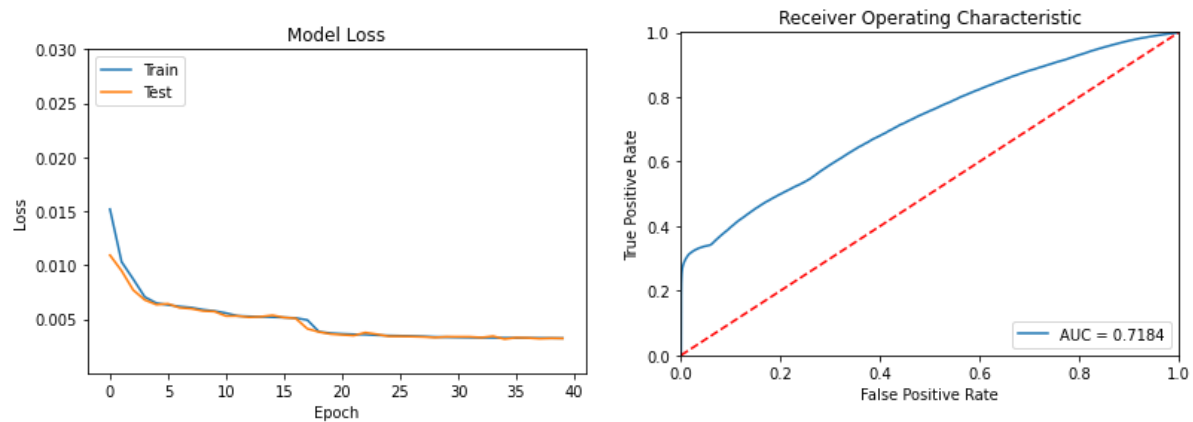


val\_loss (mse): 3.6713e-04  
val\_accuracy: 0.7190

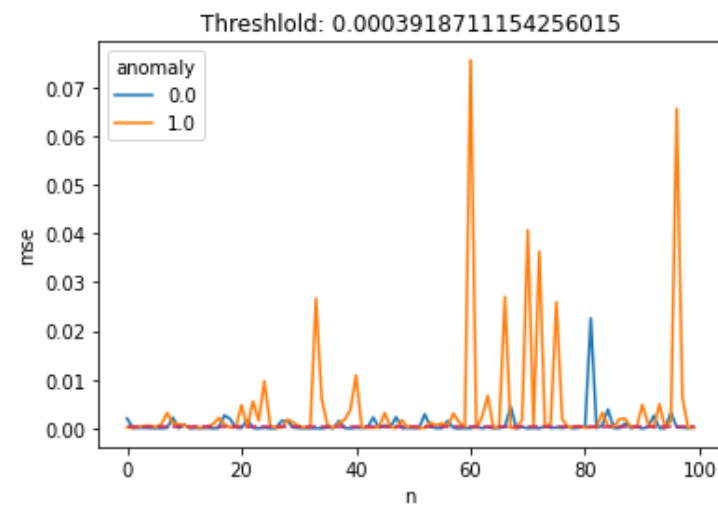
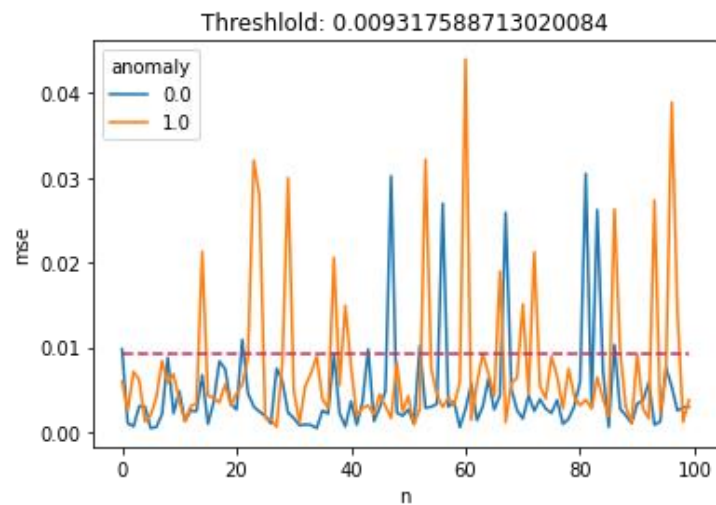
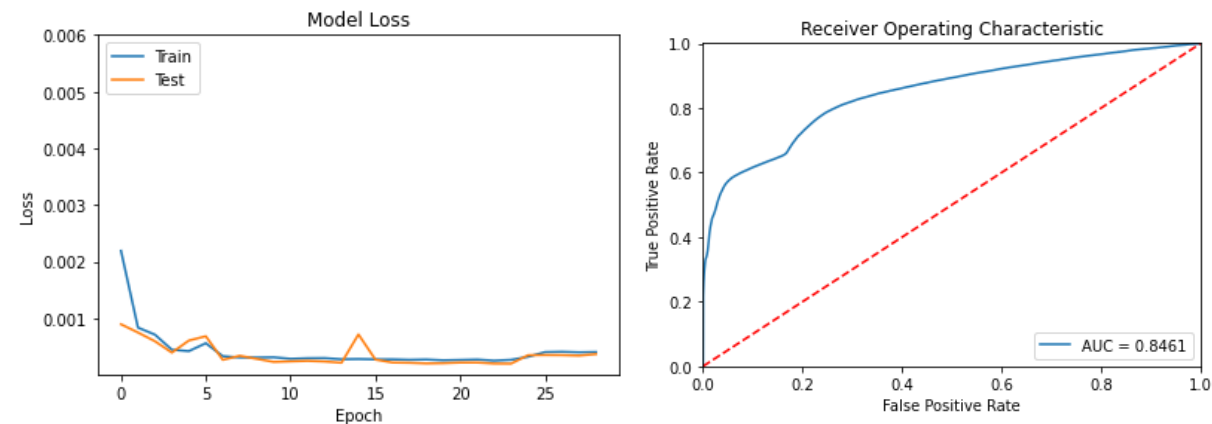
# 2 Autoencoder

## Evaluation

Basic Model



Improved Model

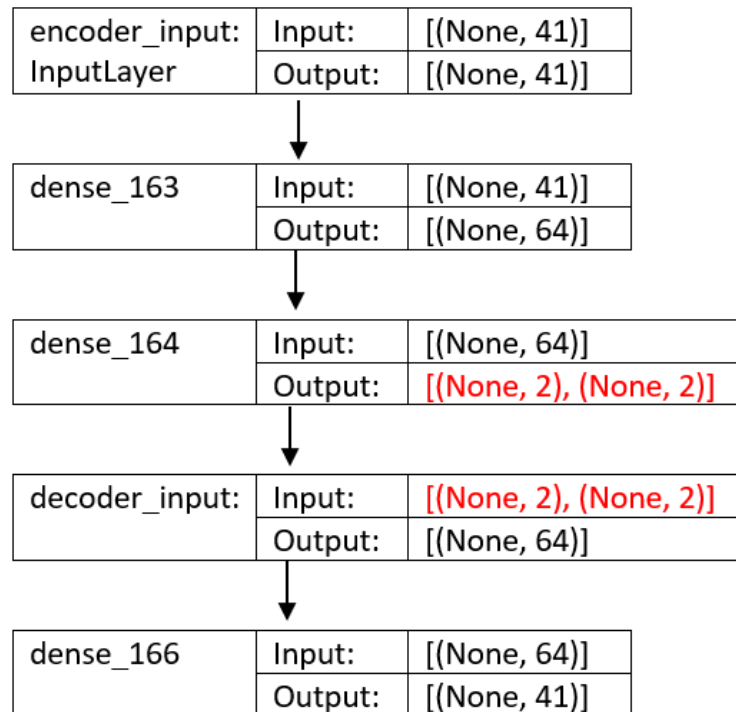


# 3 Variational Autoencoder

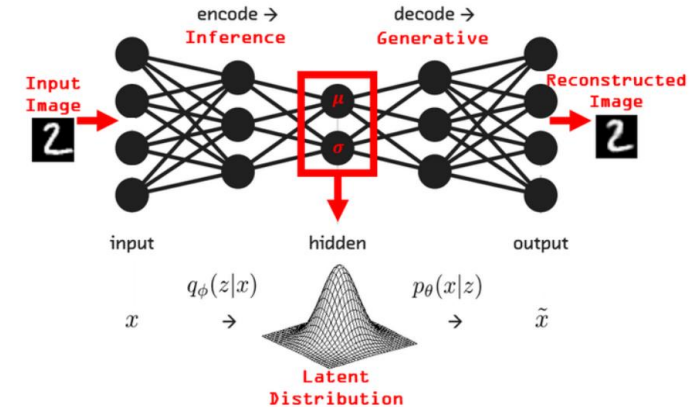
## Modeling

Parameter: loss=reconstruction\_loss + kl\_loss, optimizer='adam'

### Basic Model



val\_loss : 10.6706  
val\_accuracy: 0.3098



- Change architectures
  - Hidden layers

More complex architecture:  
Adding hidden layer functions little

- Parameters
  - Batch size
  - optimizer

Optimizer: (1) RMSprop doesn't work  
(2) Adamax is sort of better than Adam

- Early Stopping

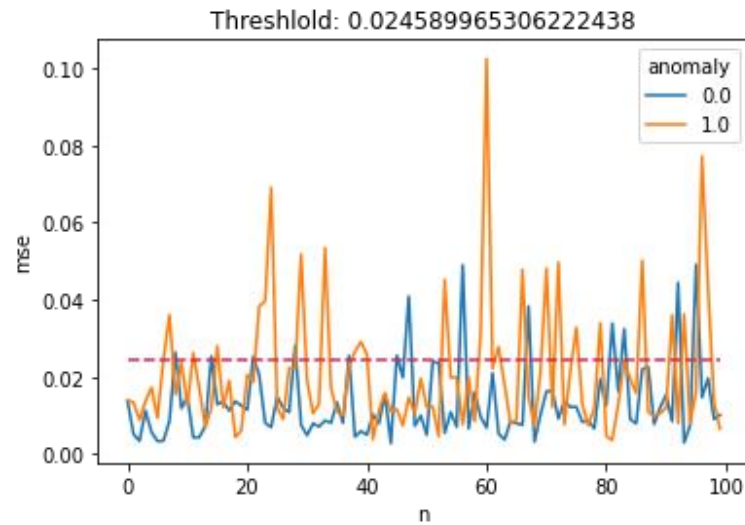
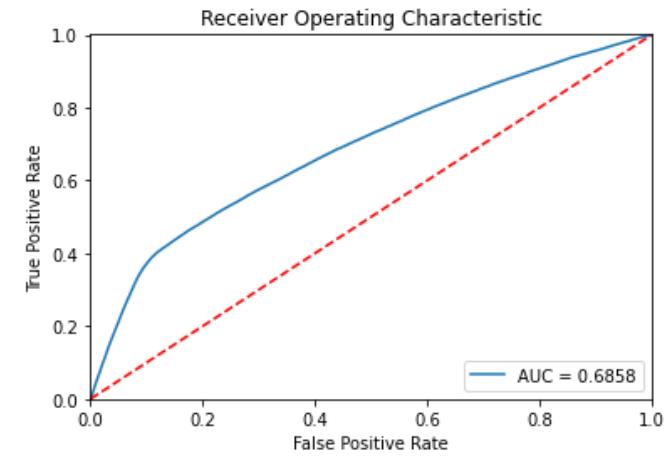
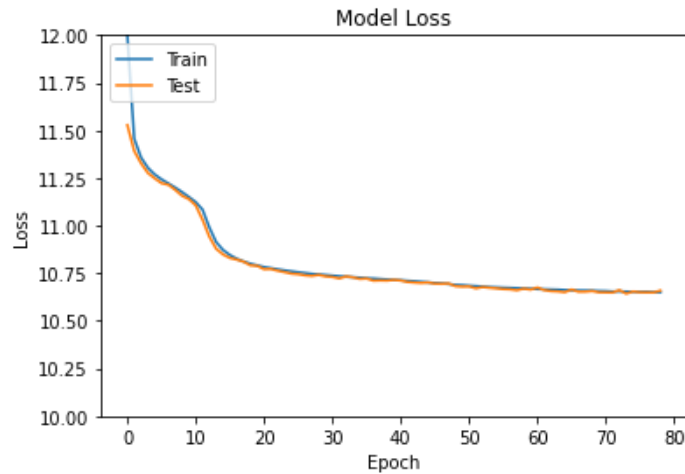
**Models improve little**



# 3 Variational Autoencoder

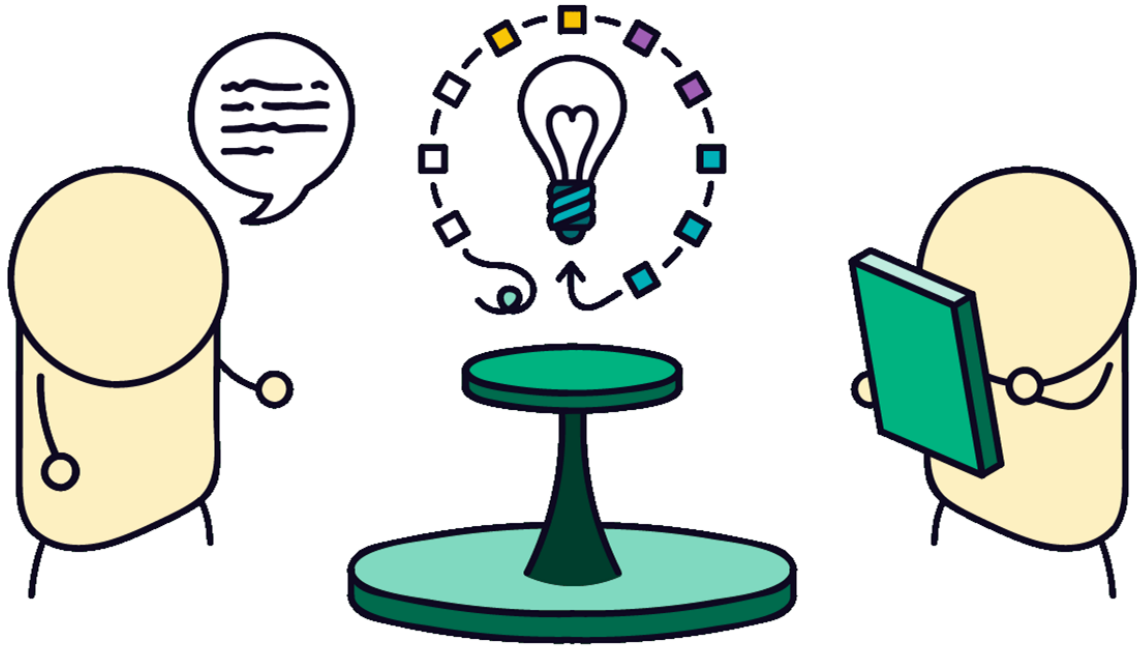
## Evaluation

### Basic Model



# 4 Conclusion

- Autoencoder can perform well in this anomaly detection problem after improving model using change architectures, parameters, and early stopping
- Potential Method:
  - (1) Use pretraining model to train the autoencoder
  - (2) Use other optimizers to train model (simulated annealing, genetic algorithm, particle swarm optimization, etc.)
- VAE doesn't work well so far
- Potential Solutions: Use other optimizers to train model

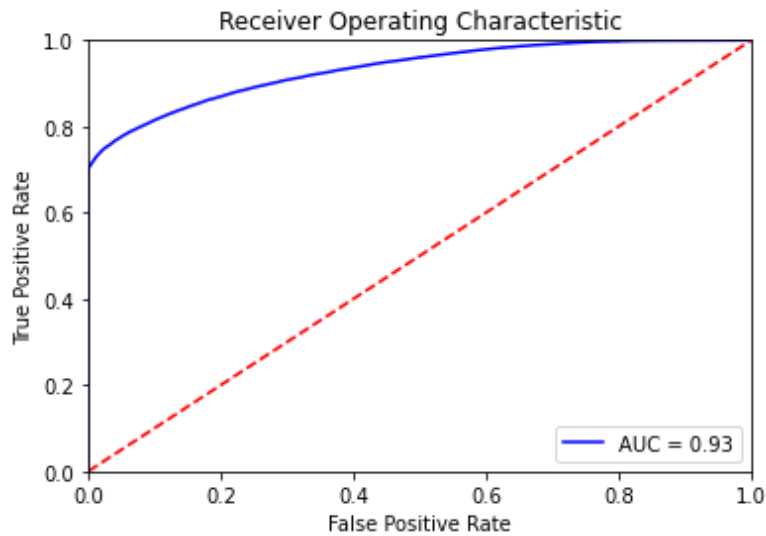


## 4 Model Performance Comparison

# Model Performance Comparison

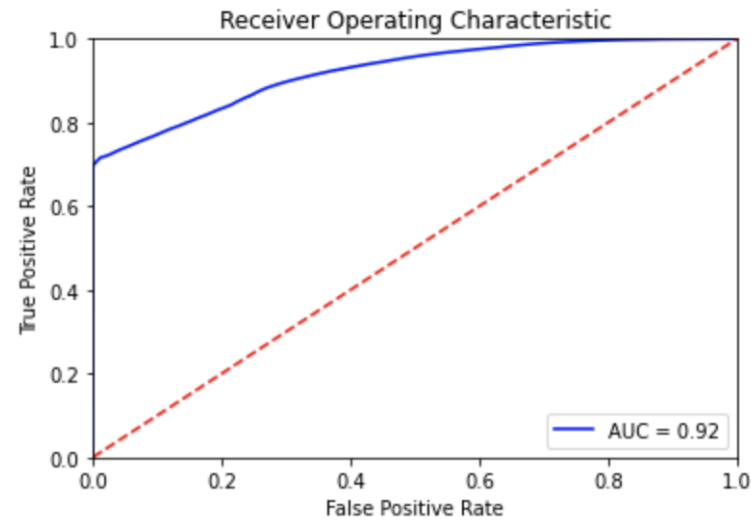
(1) Fully-connected NN-Undersampling

Accuracy: 0.8642



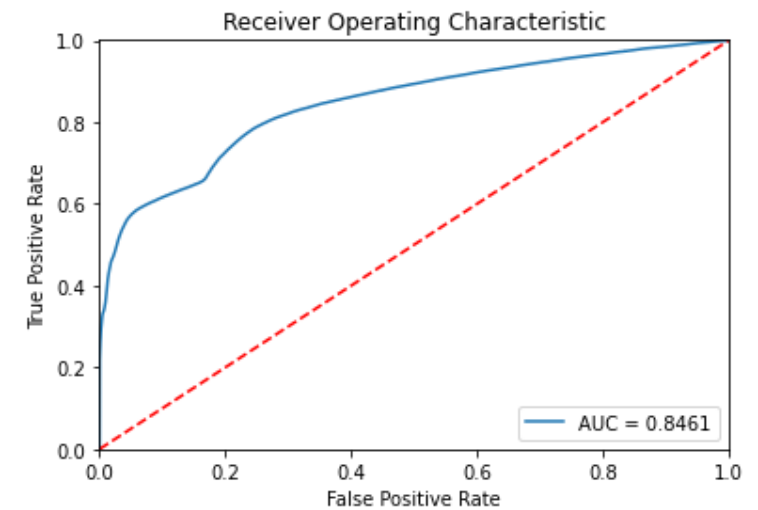
(2) CNN-Oversampling

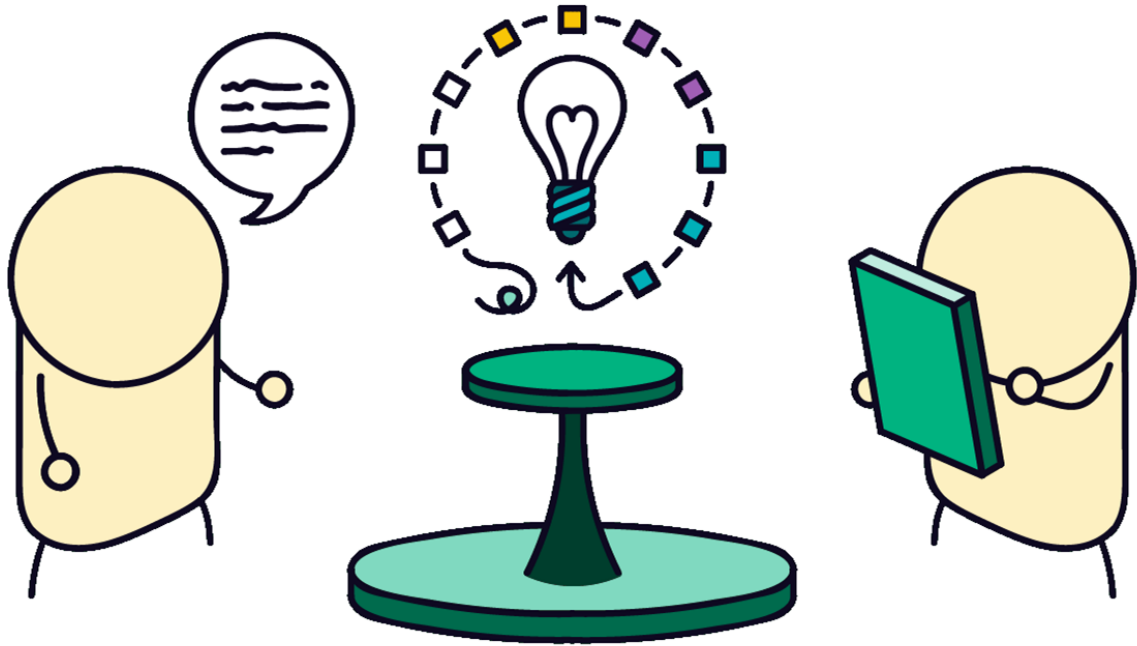
Accuracy: 0.8485



(3) Autoencoder

Accuracy: 0.6988





## 5 Conclusions

# Conclusion

## **Fully-connected NN:**

- **Undersampling** would ease imbalance problem, and functions better than the oversampling tools for fully-connected NN and this dataset

## **CNN:**

- **Oversampled** data performs better with the CNN model in both accuracy and AUC/ROC curve

## **Autoencoder/VAE:**

- Based on current algorithm, autoencoder performs better than VAE

For current deep learning algorithms and dataset, **undersampled fully-connected NN** works best for this loan default detection problem.

# Sources

- [1] Brownlee, Jason. “SMOTE for Imbalanced Classification with Python.” Machine Learning Mastery, 20 Aug. 2020, [machinelearningmastery.com/smote-oversampling-for-imbalanced-classification/](https://machinelearningmastery.com/smote-oversampling-for-imbalanced-classification/).
- [2] Berk Gokden, Applying Anomaly Detection with Autoencoders to Fraud Detection, towards data science 2020: <https://towardsdatascience.com/applying-anomaly-detection-with-autoencoders-to-fraud-detection-feaaee6b5b09>
- [3] Deep Dense Convolutional Networks for Repayment Prediction in Peer-to-Peer Lending, [http://sclab.yonsei.ac.kr/publications/Papers/IC/2018\\_SOCO\\_JYK.pdf](http://sclab.yonsei.ac.kr/publications/Papers/IC/2018_SOCO_JYK.pdf)
- [4] <https://medium.com/@RaghavPrabhu/understanding-of-convolutional-neural-network-cnn-deep-learning-99760835f148>
- [5] LendingClub Issued Loans Data: <https://www.kaggle.com/husainsb/lendingclub-issued-loans>
- [6] Renström, Holmsten: Fraud Detection on Unlabeled Data with Unsupervised Machine Learning, KTH 2018
- [7] The Keras Blog-Building Autoencoders in Keras: <https://blog.keras.io/building-autoencoders-in-keras.html>
- [8] Tom Sweers, Autoencoding Credit Card Fraud, Radboud University 2018
- [9] Understanding of Convolutional Neural Network (CNN) — Deep Learning  
Varun Chandola, Arindam Banerjee, and Vipin Kumar. Anomaly detection: A survey. ACM Comput. Surv., 41(3):15:1–15:58, July 2009
- [10] Z. Chen, C. K. Yeo, B. S. Lee, and C. T. Lau. Autoencoder-based network anomaly detection. In 2018 Wireless Telecommunications Symposium (WTS), pages 1–5, April 2018.
- [11] Packages: Python 3.0 numpy; pandas; matplotlib; seaborn; sklearn; mlxtend.plotting; pathlib2; pickle; keras; tensorflow;